

Network : Transport Protocols

Jae Hyeon Kim

— Reference

William Stalling, Data and Computer Communications 10/E, Prentice Hall

— Transport Protocols

- End-to-end transport of data that shields the user from the details of the underlying communication systems
 - Reliable, connection oriented : has greater complexity, e.g. TCP
 - Best effort, connectionless : datagram, e.g. UDP
- Connection-oriented transport protocol mechanisms
 - Provides establishment, maintenance and termination of a logical connection between TS(Transport Service) users
 - Most common service for a wide variety of applications
 - Is reliable, but complex

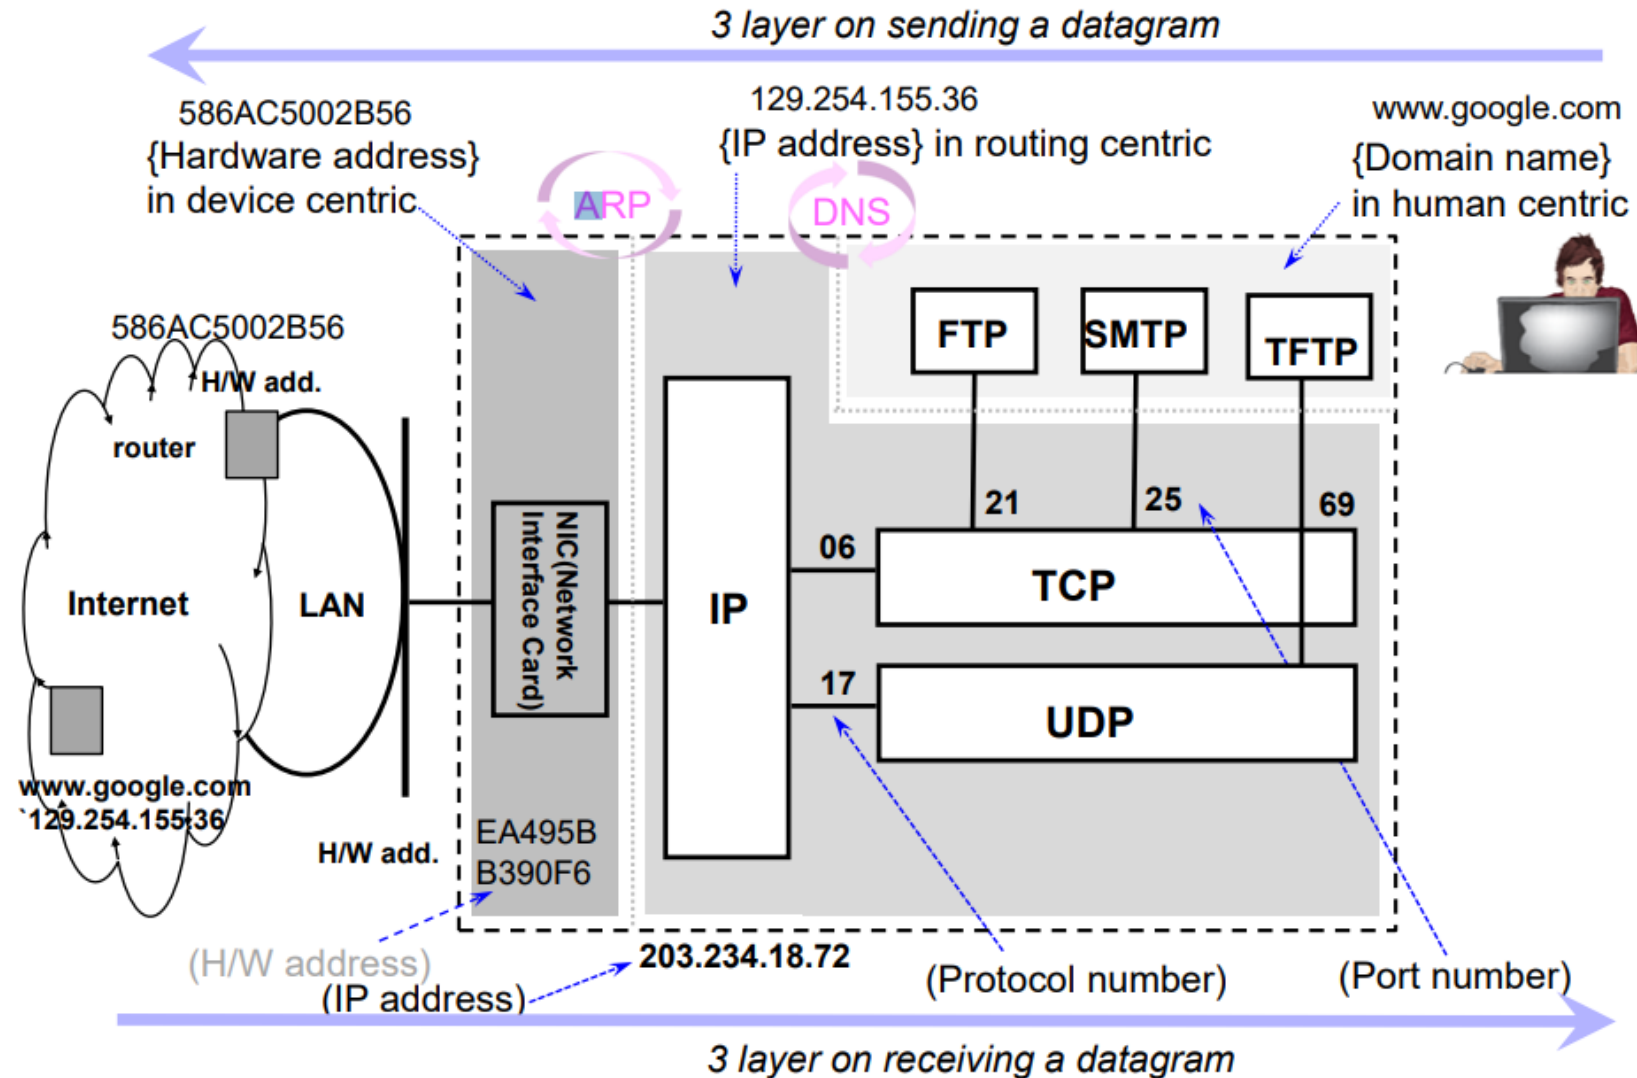
— Reliable Sequencing Network Service

- Transport service is a simple end to end protocol between two systems on the same network
- However, its situation on internet would be so complex
 - Because the below layer, so IP, provide a best-effort delivery
- Issues are :
 - Addressing
 - Multiplexing
 - Flow control
 - Connection establishment and termination

— Addressing


- Transport protocol must be able to derive the following information from the TS user address
- User identification
 - A socket in TCP
 - Port represents a particular transport service user
 - Multiple users employ the same transport protocol and are distinguished by port numbers
- Transport entity identification (on host)
 - Specify transport protocol (TCP, UDP)
- Host address of attached network device
 - In the internet, a global internet address

3-layer Internet Addressing



cf) Domain Name

- Problem statement
 - Average brain can easily remember 7 digits
 - On average, IP addresses have 10.28 digits
- So, we need an easier way to remember IP addresses
 - Makes use of alphanumeric names to refer to hosts
 - Add a distributed, hierarchical protocol (called DNS)
 - Address resolution is to map between host names and IP add.
- Let's re-think our DN
 - www.chonbuk.ac.kr vs www.jbnu.ac.kr




— Flow Control

- Complex at the transport layer:
 - Considerable delay in the comm. Of flow control information
 - The transmission delay may be highly variable, making it difficult to effectively use a timeout mechanism for retransmission of lost data
- Reasons for control:
 - User of the receiving transport entity cannot keep up with the flow
 - Receiving transport entity itself cannot keep up with the flow of segments

— Alternative to Flow Control Requirements

- Do nothing
 - Segments that overflow are discarded
 - Sending transport entity will fail to get ACK and will retransmit
 - Thus further adding to incoming data
- Refuse to accept further segments from the network service
 - Relies on network service to do the work
 - clumsy
- Use a fixed sliding widow protocol
 - With a reliable network service this works quite well
- Use a credit scheme
 - A more effective scheme to use with an unreliable network service

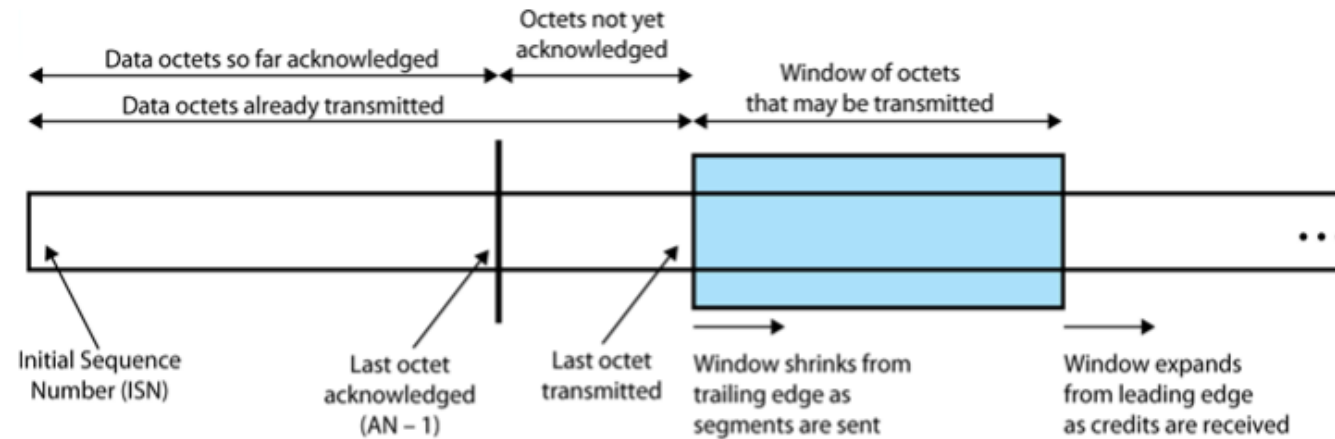
— An Example of the Fixed Sliding Window



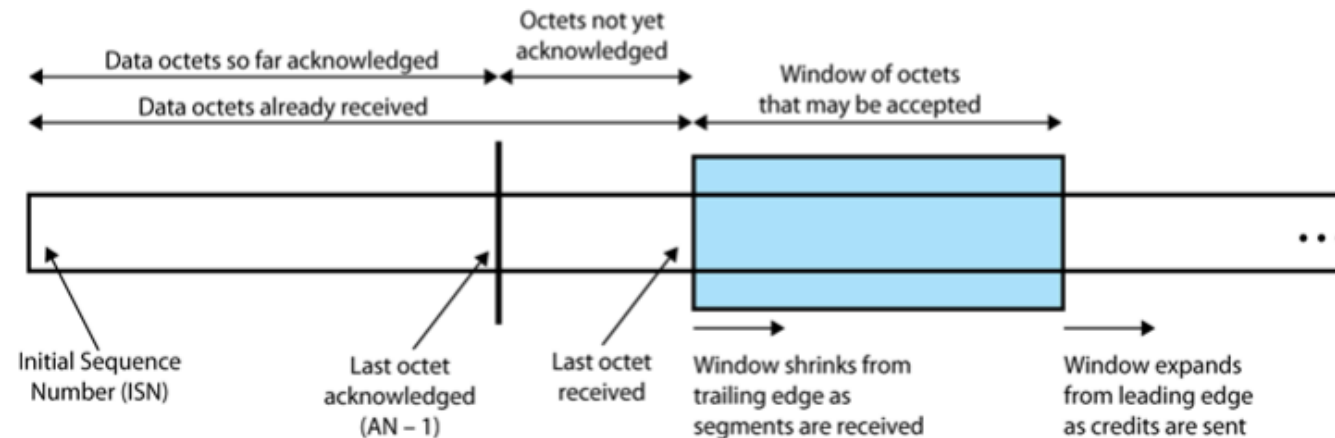
— Credit Scheme

- Decouples flow control from ACK
 - Then, credit scheme mainly rules the flow control
- Each octet has sequence number
- Each transport segment has sequence number (SN), Ack. Number (AN) and window size (W) in header
- Sends sequence number of first octet in segment
- ACK includes (AN= I , W= j) which means
 - All octets through SN= $i-1$ acknowledged, want I next
 - Permission to send additional window of W= j octets

Sending and Receiving Perspectives

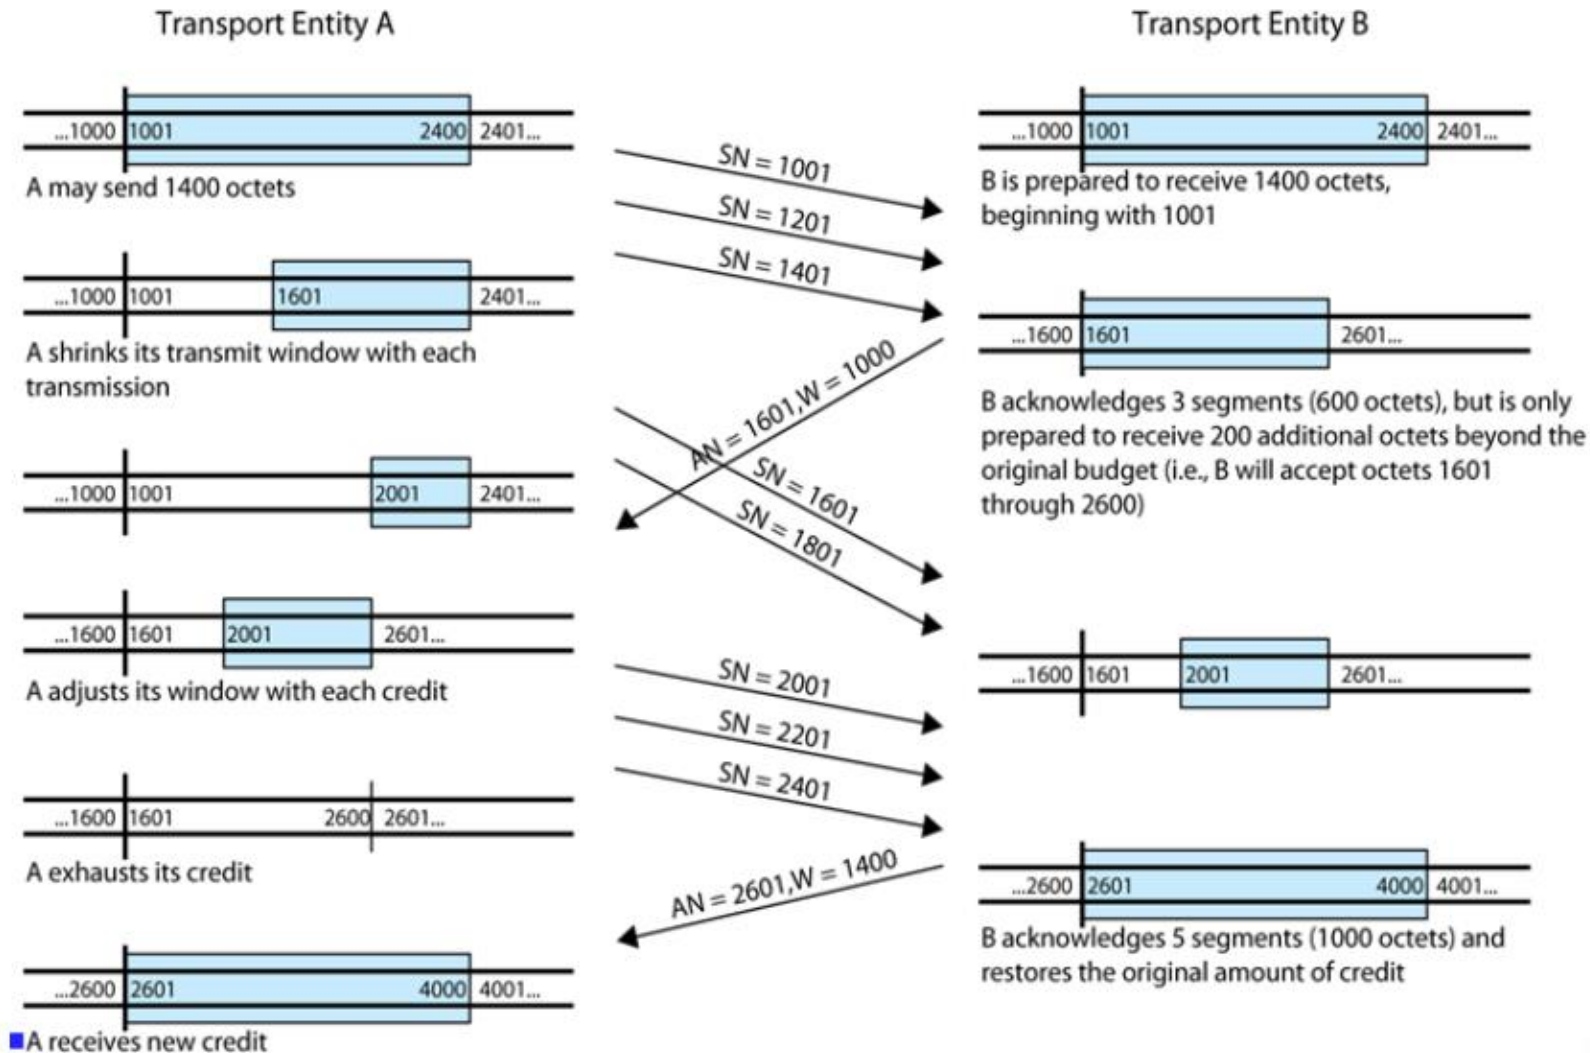


(a) Send sequence space



(b) Receive sequence space

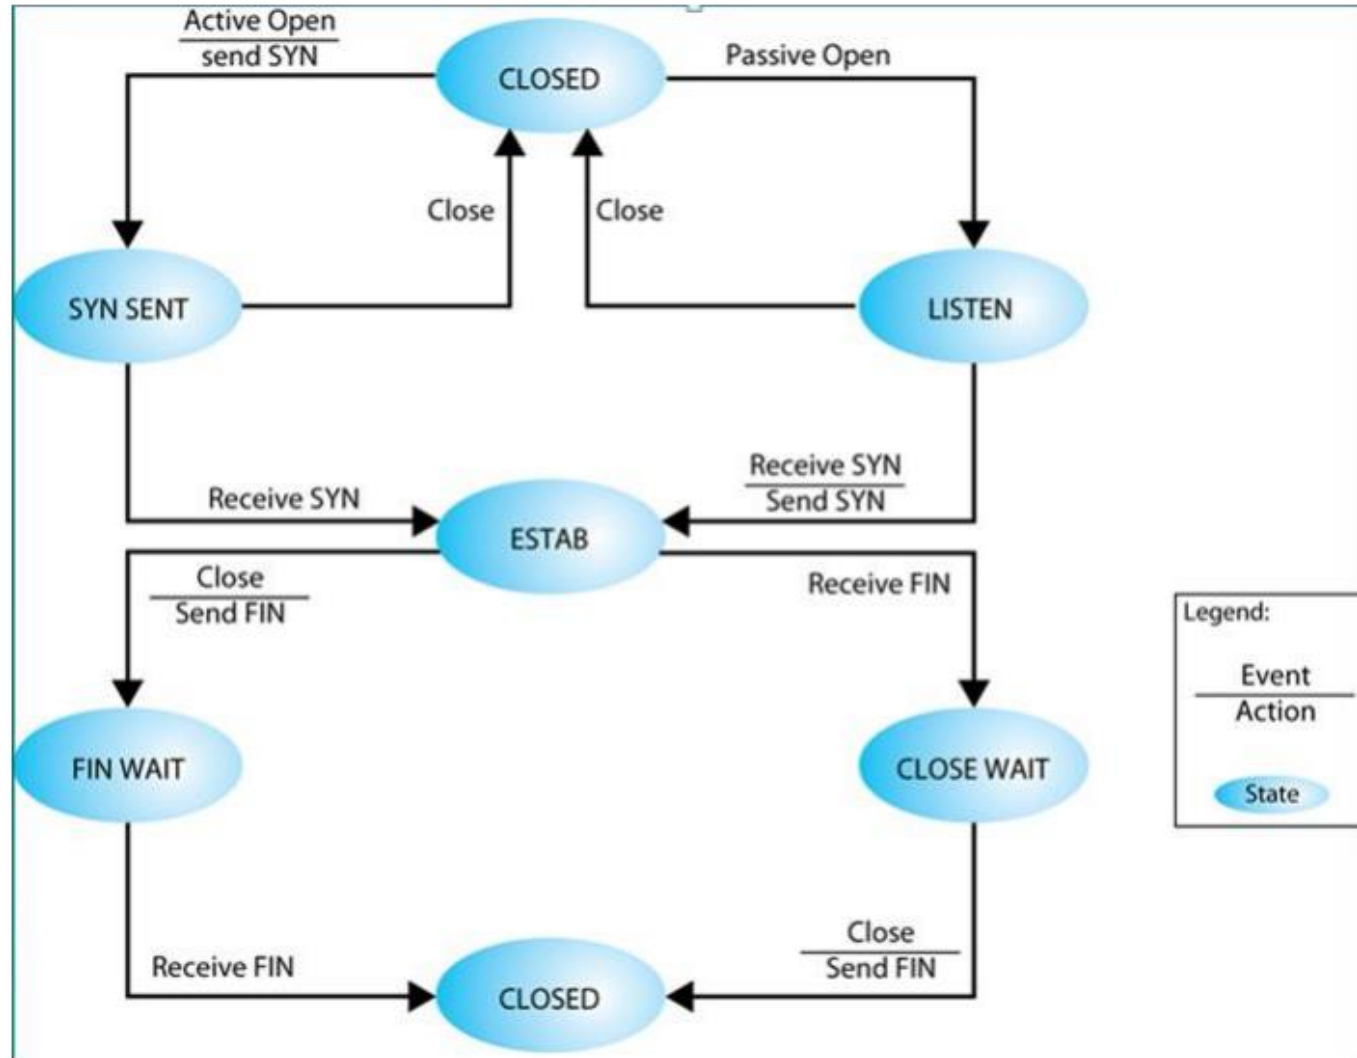
Credit Allocation Mechanism



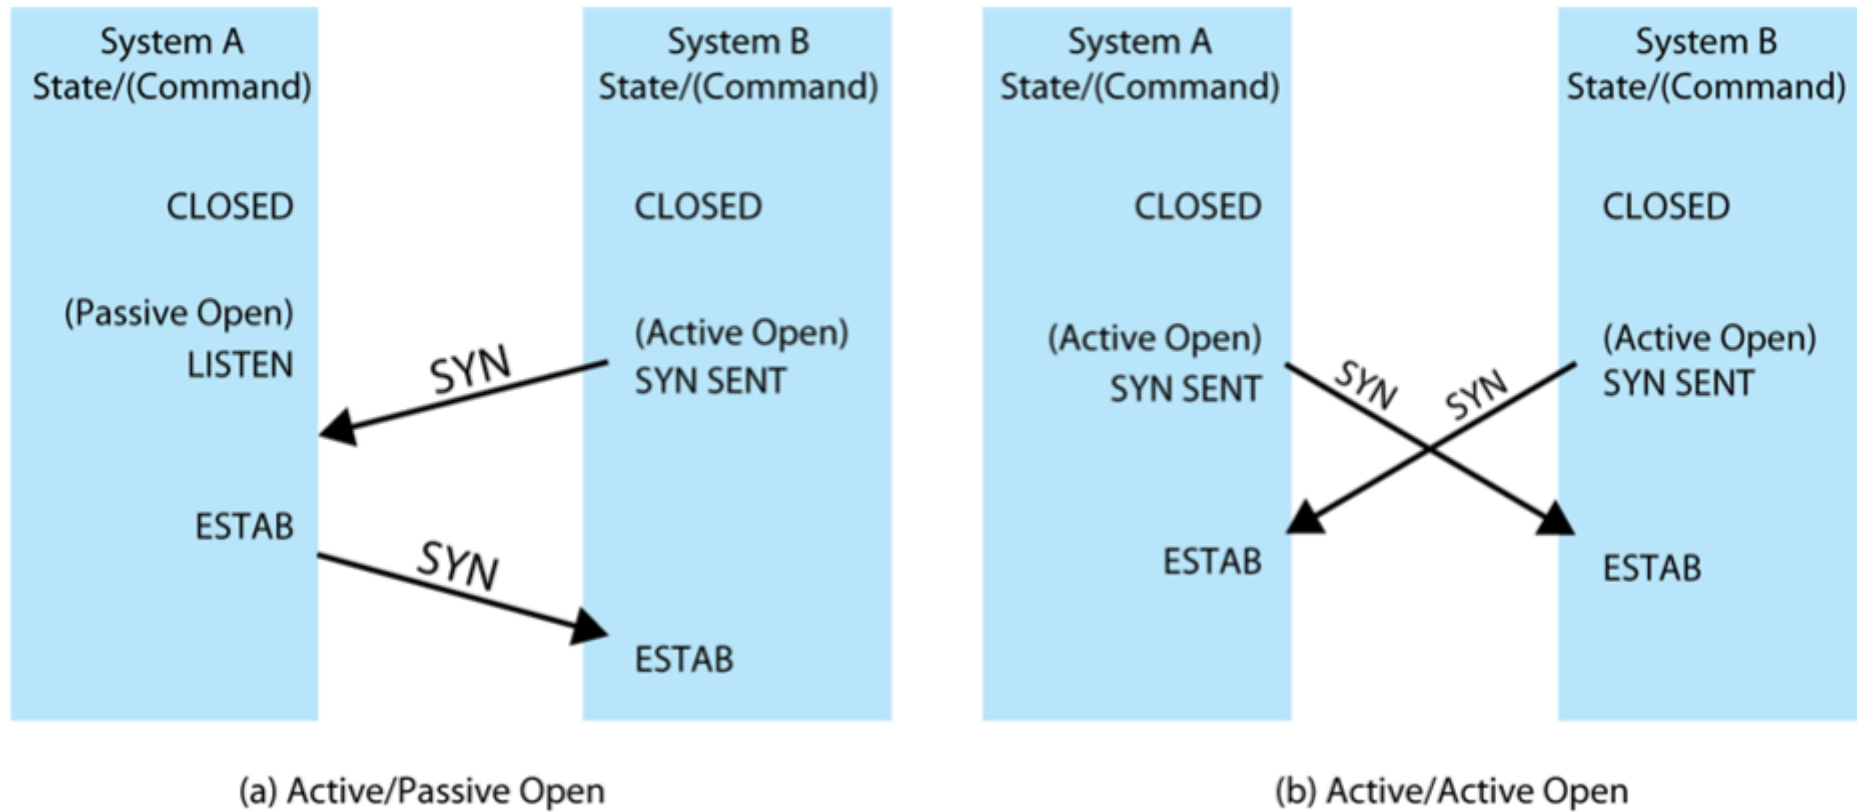
— Connection Establishment and Termination

- Serves three main purpose:
 - Allow each end to assure that the other exists, and the other's intension to communicate with the initiator
 - Allows exchange or negotiation of optional parameters
 - Triggers allocation of transport entity resources
- By mutual agreement

– Simple Connection State Diagram



— Connection Establishment Scenario



— Unreliable Network Service

- More difficult case for transport protocol since
 - Segments may get lost and/or arrive out of order
- Examples include
 - Internetwork using IP
 - IEEE 802.3 with unacknowledged connectionless LLC service
- Issues:
 - Ordered delivery
 - Retransmission strategy
 - Duplication detection
 - Flow control
 - Connection establishment & termination
 - Crash recovery

— Ordered Delivery

- With an unreliable network service it is possible that segments may arrive out of order
- Solution is to number segments sequentially
 - TCP makes use of a scheme where each data octet is implicitly numbered


— Retransmission Strategy

- Events necessitating retransmission:
 - Segment may be damaged in transit but still arrives at its dest.
 - Segment fails to arrive
- Sender does not know transmission was unsuccessful
- Receiver acknowledges successful receipt by returning a segment containing an acknowledgement number
- Retransmission strategy
 - No ACK will be issued if a segment does not arrive successfully
 - A timer needs to be associated with segment at it is sent
 - If timer expires before acknowledgement is received, sender must retransmit

— Duplication Detection

- Receiver received a segment, but lost its Ack.
 - Eventually the sender try to retransmit the segment
 - Receiver must be able to recognize duplicates
 - Segment sequence numbers help
- Complications arise if:
 - A duplicate is received prior to the close of the connection
 - Sender must not get confused if it receives multiple acknowledgements to the same segment
 - Sequence number space must be long enough
 - A duplicate is received after the close of the connection

Incorrect Duplicate Detection




— Flow Control

- Credit allocation quite robust with unreliable network
 - Can ack data & grant credit
 - Lost ACK recovers on next received
- Have problem if $AN=1$, $W=0$ closing window
 - Then send $AN=1$, $W=j$ to reopen, but this is lost
 - Sender thinks window closed, receiver thinks it open
- Solution is to use persist timer
- If timer expires, send something
 - Could be re-transmission of previous segment


— Connection Establishment

- Two way handshake
 - A send SYN, B replies with SYN
 - Lost SYN handled by re-transmission
 - Ignore duplicate SYNs once connected
- Lost or delayed data segments can cause connection problems
 - With data segment from old connections
 - Make use SYN i , where i is the sequence # of the first data
 - Start segment # far removed from previous connection
 - There is still a problem with SYN segment from old connections
 - Need ACK to include i
 - So, three way handshake


Two Way Handshake problem with Obsolete Data Segment



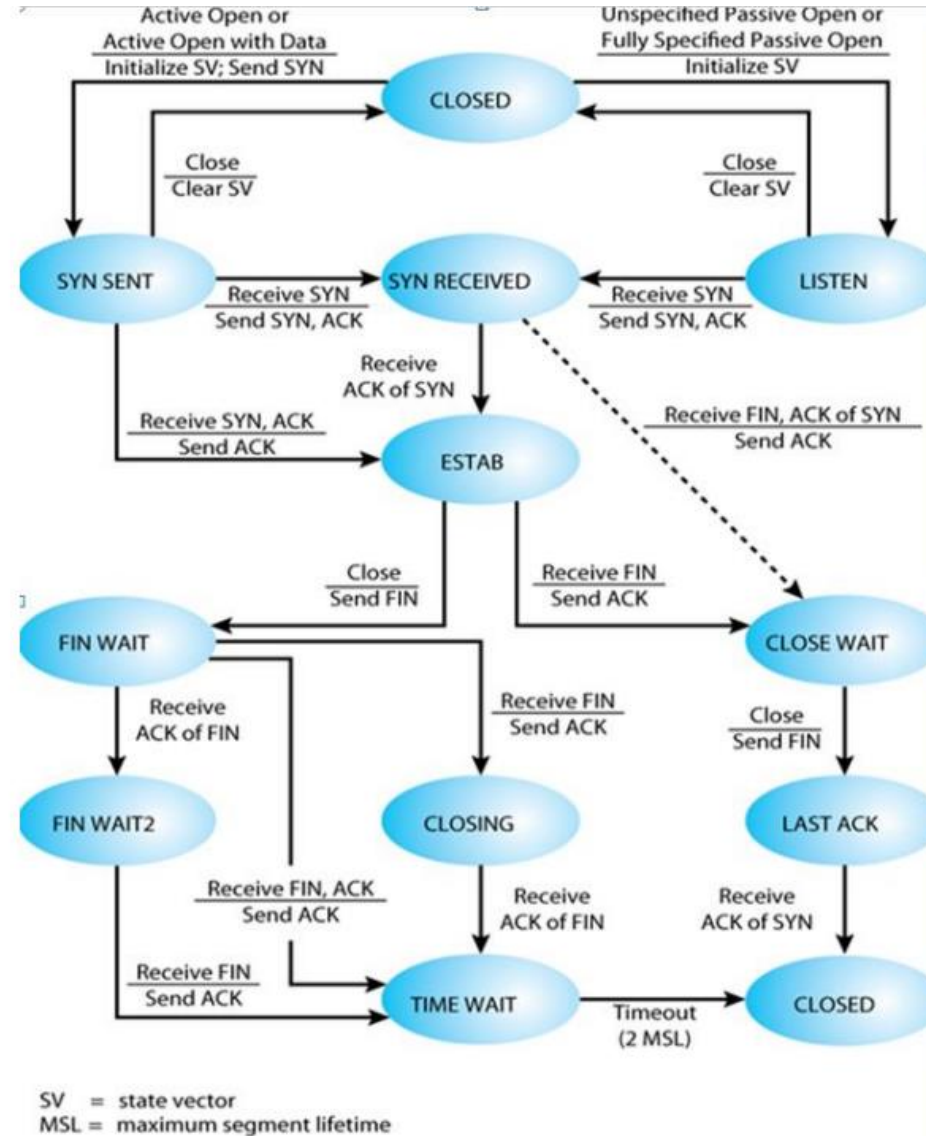
— Obsolete SYN Segment



Three Way Handshake : Examples



TCP Entity State Diagram



— Connection Termination

- 2-way handshake is inadequate on an unreliable network
 - Like connection establishment, it needs 3-way handshake
- Out of order segments could cause:
 - The FIN segment to arrive before the last data segment
- To avoid this problem the next sequence number after the last octet of data can be assigned to FIN
 - Each side must explicitly acknowledge the FIN of the other using an ACK with the sequence number of the FIN to be acknowledged

— Failure Recovery

- After restart a system, the state information of all active connections is lost
 - May have half open connection because side that did not crash still thinks it is connected
- Still active side of a half-open connection can close the connection using a keepalive timer
 - Wait for ACK for (time out) * (number of retries)
 - When expired, close connection and inform user
- Or, failed side returns an RST i to every segment i that it receives
 - RST i must be checked for validity on the other side
 - If valid, an abnormal termination occurs

— TCP Services


- TCP (Transmission Control Protocol) : RFC 793
 - Connection oriented, reliable communication over reliable and unreliable (inter)networks
- Two ways of labeling data:
 - Data stream push
 - User requires transmission of all data up to push flag
 - Receiver will deliver in same manner
 - Avoids waiting for full buffers
 - Urgent data signal
 - Indicates urgent data is upcoming in stream
 - User decides how to handle it

TCP Services Request Primitives

Primitive	Parameters	Description
Unspecified Passive Open	source-port, [timeout], [timeout-action], [precedence], [security-range]	Listen for connection attempt at specified security and precedence from any remote destination.
Fully Specified Passive Open	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security-range]	Listen for connection attempt at specified security and precedence from specified destination.
Active Open	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security]	Request connection at a particular security and precedence to a specified destination.
Active Open with Data	source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security], data, data-length, PUSH-flag, URGENT-flag	Request connection at a particular security and precedence to a specified destination and transmit data with the request.
Send	local-connection-name, data, data-length, PUSH-flag, URGENT-flag, [timeout], [timeout-action]	Transfer data across named connection.
Allocate	local-connection-name, data-length	Issue incremental allocation for receive data to TCP.
Close	local-connection-name	Close connection gracefully.
Abort	local-connection-name	Close connection abruptly.
Status	local-connection-name	Query connection status.

Note: Square brackets indicate optional parameters.

TCP Header Format



— TCP Mechanisms

- Can be grouped into:
- connection establishment
 - Always uses a three-way handshake
 - Connection is determined by host and port
- Data transfer
 - Viewed logically as consisting of a stream of octets
 - Flow control is exercised using credit allocation
- Connection termination
 - Each TCP user must issue a CLOSE primitive
 - An abrupt termination occurs if the issues an ABORT primitive

— TCP Implementation Policy Options (1)

- Send policy
 - If no pushed data, a sending TCP entity transmits at its own convenience in credit allocation
 - May construct segment per batch of data from user : quick response but higher overheads
 - May wait for certain amount of data : slower response but lower overheads
- Deliver policy
 - In absence of push, can deliver data at own convenience
 - May deliver from each segment received : higher O/S overheads but more responsive
 - May buffer data from multiple segments : less O/S overheads but slower

— TCP Implementation Policy Options (2)

- Accept policy
- If segments arrive out of order the receiving TCP entity has two options:
- In-order
 - Accept only segments that arrive in order; any segment that arrives out of order is discarded
 - Makes for simple implementation but places a burden on the networking facility
 - If a single segment is lost in transit, then all subsequent segments must be retransmitted
- In-window
 - Accept all segments that are within the receive window
 - Requires a more complex acceptance test and a more sophisticated data storage scheme

— TCP Implementation Policy Options (3)

- Retransmit policy (three strategies)
- First only
 - Maintain one retransmission timer for entire queue
 - If timer expires, retransmit the segment at the front of the queue
 - Efficient traffic generation, but can have considerable delays
- Batch
 - Maintain one retransmission timer for entire queue
 - If timer expires, retransmit all segments in the queue
 - May result in unnecessary retransmissions

— TCP Implementation Policy Options (3)

- Individual
 - Maintain one timer for each segment in the queue
 - More complex implementation

— TCP Implementation Policy Options (4)

- Acknowledge policy (timing of ack.)
- Immediate
 - Immediately transmit an empty segment containing the appropriate acknowledgment number
 - Simple and keeps the remote TCP fully informed
 - Limits unnecessary retransmissions
 - Can cause a further load on the network
- Cumulative
 - Wait for an outbound segment with data on which to piggyback the acknowledgement
 - Typically used
 - Requires more processing at the receiving end and complicates the task of estimating round-trip time

UDP

- Connectionless service specified in RFC 768
 - Delivery & duplication control not guaranteed
- Reduced overhead in transmission
- Least common denominator service
- Uses:
 - Data collection & dissemination
 - Request-response
 - Real time application

